Exercise Set 2 - Reinforcement Learning
Chapter 3,4 - Tabular Methods

Instructions

This is the second exercise booklet for Reinforcement Learning. It covers both ungraded exercises to practice at home or during the tutorial sessions as well as graded homework exercises and graded coding assignments. The graded assignments are clearly marked.

- Make sure you deliver answers in a clear and structured format. LaTeXhas our preference. Messy handwritten answers will not be graded.

- Pre-pend the name of your TA to the file name you hand in and remember to put your name and student ID on the submission;

- The deadline for this first assignment is **September 23rd 2022 at 13:00** and will cover the material of chapter 3-4. All questions marked 'Homework' in this booklet need to be handed in on Canvas. The coding assignments need to be handed in separately through the Codegra.de platform integrated on canvas.

Contents

Chapter 3: Monte Carlo methods

3.1   Monte Carlo

1. Consider an MDP with a single state s_0 that has a certain probability of transitioning back onto itself with a reward of 0, and will otherwise terminate with a reward of 5. Your agent has interacted with the environment and has gotten the following two trajectories: $[0, 0, 5], [0, 0, 0, 5]$. Use $\gamma = 0.9$.

   (a) Estimate the value of s_0 using first-visit MC.

   (b) Estimate the value of s_0 using every-visit MC.

3.2   Importance Sampling in Monte Carlo methods

1. What is a disadvantage of using *ordinary importance sampling* in off-policy Monte Carlo?

2. What is a disadvantage of using *weighted importance sampling* in off-policy Monte Carlo?

3.3   Bias of v_π Monte Carlo estimators

Consider one episode following π: $(S_0, A_0, R_1, S_1, A_1, R_2, \ldots, S_{T-1}, A_{T_1}, R_T)$, where $S_0 = s$. Determine and provide intuition on the biasedness of the following estimators for $v_\pi(s)$:

1. $\sum_{i=1}^{T} \gamma^{i-1} R_i$.

2. $\frac{1}{|J|} \sum_{j \in J} \sum_{i=1}^{T-j} \gamma^{i-1} R_{j+i}$, where J contains all indices j such that $S_j = s$.

3. $\sum_{i=1}^{T-t_s} \gamma^{i-1} R_{t_s+i}$, where t_s is the latest time step such that $S_{t_s} = s$. How does this compare to the first-visit MC estimator?

3.4   * Exam question: Monte Carlo for control

This exercise has been taken from a previous exam (perhaps lightly edited) and can require a bit more insight. It should give you a good idea of the level of exam questions.
   The following questions refer to the pseudo-code presented in Figure 1.

1. Part of the algorithm is covered by a black square. What is the missing information?

2. Is this a Monte-Carlo algorithm or a TD-based algorithm? Explain your answer based on the given pseudo-code.

3. What is stored in $C(S_t, A_t)$?

4. Why is the inner loop stopped when $A_t \neq \pi(S_t)$?

3.5   Homework: Coding Assignment - Monte Carlo

1. In the chapter 5.6 of the book, we are given incremental update rule for weighted importance sampling (equations 5.7-5.8). However, in the coding assignment we will use ordinary importance sampling which uses a different update rule. Start with $V_n = \frac{\sum_{k=1}^{n} W_k G_k}{n}$ and derive the incremental update rule for ordinary importance sampling. Your final answer should be of the form: $V_n = V_{n-1} + a * (b - V_{n-1})$.

Figure 1: Algorithm pseudo-code.

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:
    $Q(s, a) \in \mathbb{R}$ (arbitrarily)
    $C(s, a) \leftarrow 0$
    $\pi(s) \leftarrow \arg\max_a Q(s, a)$ (with ties broken consistently)

Loop forever (for each episode):
    $b \leftarrow$ any soft policy
    Generate an episode using b: $S_0, A_0, R_1, \ldots, S_{T-1}, A_{T-1}, R_T$
    $G \leftarrow 0$
    $W \leftarrow 1$
    Loop for each step of episode, $t =$ ▮▮▮▮▮▮▮▮▮:
        $G \leftarrow \gamma G + R_{t+1}$
        $C(S_t, A_t) \leftarrow C(S_t, A_t) + W$
        $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$
        $\pi(S_t) \leftarrow \arg\max_a Q(S_t, a)$ (with ties broken consistently)
        If $A_t \neq \pi(S_t)$ then exit inner Loop (proceed to next episode)
        $W \leftarrow W \frac{1}{b(A_t | S_t)}$

2. Download the notebook *RLLab2_MC.zip* from canvas assignments and follow the instructions.

3. What is the difference between Dynamic Programming and Monte Carlo? When would you use the one or the other algorithm? Include your answer in the submitted homework.

Figure 2: Example MDP

Chapter 4: Temporal Difference Learning

4.1 Temporal Difference Learning (application)

Consider an undiscounted Markov Decision Process (MDP) with two states A and B, each with two possible actions 1 and 2, and a terminal state T with $V(T) = 0$. The transition and reward functions are unknown, but you have observed the following episode using a random policy:

- $A \xrightarrow[r_3=-3]{a_3=1} B \xrightarrow[r_4=4]{a_4=1} A \xrightarrow[r_5=-4]{a_5=2} A \xrightarrow[r_6=-3]{a_6=1} T$

where the the arrow ($\rightarrow$) indicates a transition and a_t and r_t take the values of the observed actions and rewards respectively.

1. What are the state(-action) value estimates $V(s)$ (or $Q(s,a)$) after observing the sample episode when applying
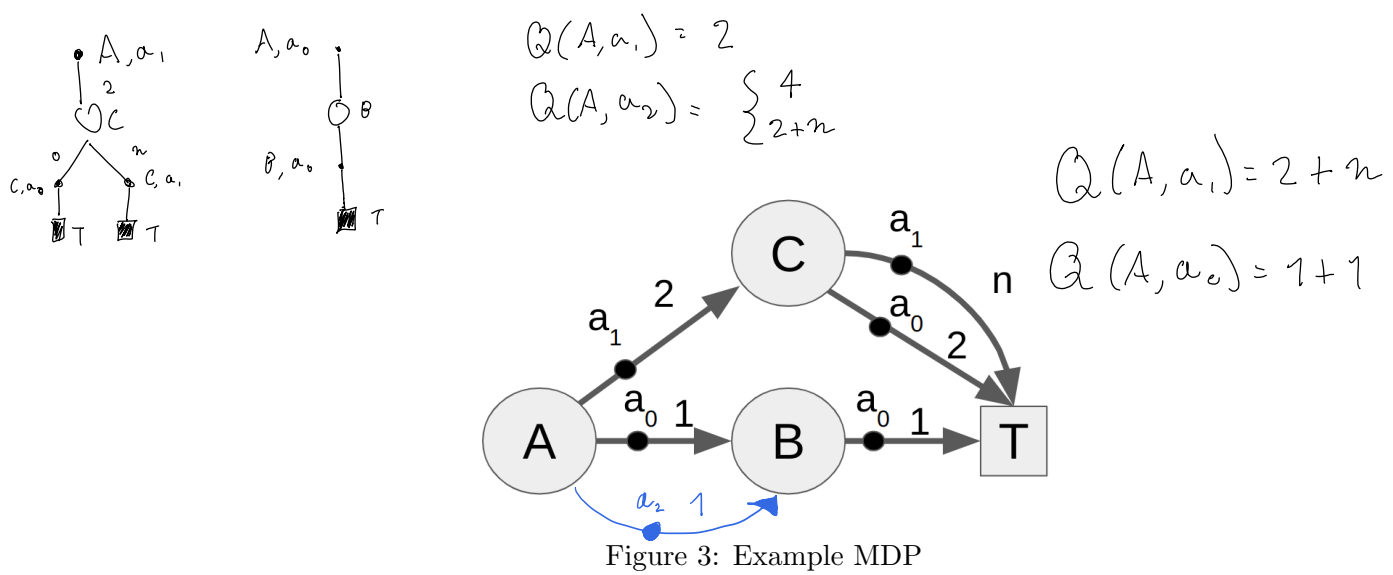
   (a) TD(0) (1-step TD)
   (b) SARSA,

   where we initialize state(-action) values to 0 and use a learning rate $\alpha = 0.1$.

4.2 Temporal Difference Learning (Theory)

1. We can use Monte Carlo to get value estimates of a state with $V_M(S) = \frac{1}{M}\sum_{n=1}^{M} G_n(S)$ where $V_M(S)$ is the value estimate of state S after M visits of the state and $G_n(S)$ the return of an episode starting from S. Show that $V_M(S)$ can be written as the update rule $V_M(S) = V_{M-1}(S) + \alpha_M[G_M(S) - V_{M-1}(S)]$ and identify the learning rate α_M.

2. Consider the TD-error
$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t). \tag{1}$$

   (a) What is $\mathbb{E}[\delta_t|S_t = s]$ if δ_t uses the true state-value function V^π
   (b) What is $\mathbb{E}[\delta_t|S_t = s, A_t = a]$ if δ_t uses the true state-value function V^π

$Q(A, a_1) = 2$

$Q(A, a_2) = \begin{cases} 4 \\ 2+n \end{cases}$

$Q(A, a_1) = 2 + n$

n $Q(A, a_2) = 1 + 1$

Figure 3: Example MDP

4.3 Off-policy TD

Consider the MDP in Figure 2. Consider a uniform behavior policy b (probability of a_1 and a_2 is 0.5 in both states). Additionally, consider the target policy π which takes a_1 with probability 0.1 and a_2 with probability 0.9 in both states. We consider the undiscounted case ($\gamma = 1$).

1. What are the Q functions Q^b and Q^π under both policies?

2. Now consider a dataset gathered using b $(s_1, a_2, 0, s_2, a_1, -1), (s_1, a_2, 0, s_2, a_2, +1)$. Consider a Q-function that is initialized as per the following table

	a_1	a_2
s_1	-1	0.5
s_2	-1	+1

   Apply one pass of Sarsa on the dataset with a learning rate of 0.1. How does the change in Q function relate to the two functions you calculated in sub-question 1?
   hint: throughout this question, only $Q(s_1, a_2)$ will change. Why?

3. We can use expected Sarsa to estimate the Q-function Q^π. Apply a single pass, and note how the change in Q function relates to the two functions of sub-question 1.

4. Another possibility for off-policy learning is applying Sarsa with importance weight. Again do one pass and notice the change in Q-function.

5. We could also learn a V function, e.g. through TD(0), off policy. For example, by using importance weights. Why do you think the book doesn't cover this?

6. Could you do something like expected Sarsa for learning a V function? If yes, apply one pass. If not, explain why this is the case.

7. Why is Q-learning considered an off-policy control method? (ex. 6.11 in RL:AI)

4.4 Homework: SARSA and Q-learning

Consider the MDP in Figure 3 with states A (starting state), B, C, T (terminal state) and rewards for corresponding actions indicated in the diagram. The reward for action a_1 in state C is parametrized by n. We use an ϵ-greedy policy with $\epsilon = 0.1$ for SARSA and the same policy as a behavior policy for Q-learning. We do not use discounting ($\gamma = 1$).

1. Suppose that $n \in (-\infty, 2)$ and we run SARSA until convergence.

(a) What would be the final Q-values for all states and actions? When necessary write your answer using an expression that includes n.

(b) For which range of values of n will the final policy prefer to choose action a_1 more often in state A? For which range of values will it prefer to choose a_0? Use final Q-values to explain your answers.

2. Suppose that $n \in (-\infty, \infty)$ and we run Q-learning until convergence.

(a) What would be the final Q-values for all states and actions? When necessary write your answer using an expression that includes n.

(b) For which range of values of n will the final policy prefer to choose action a_1 more often in state A? For which range of values will it prefer to choose a_0? Use final Q-values to explain your answers.

3. Suppose that $n = 1$ and we run both SARSA and Q-learning until convergence and record the final performance (average return after convergence). Then we use an adjusted MDP, where there is an extra action a_2 in A, that moves to B and gives reward 1. We also run both algorithms until convergence and record the final performance. Would the final performance (average return after convergence) of SARSA or Q-learning be different across these two MDPs? Explain your answer.

Note: You don't need to compute the actual values for final performance to answer this question. Make an argument that explains/supports your answer.