Exercise Set 3 - Reinforcement Learning

Advanced TD methods and approximation

Giulio Starace - 13010840

September 30, 2022

Homework: Coding Assignment - Temporal Difference Learning

1. Coding answers have been submitted on codegra under the group "stalwart cocky sawly".

2. As can be seen in Fig. 1, Q-learning achieves higher average return for the particular environment considered. This is the opposite of Example 6.6 in the book where SARSA achieves higher average return. This is due to the difference between environments. In Example 6.6, the environment contains a penalty cliff adjacent to the optimal path, causing optimally acting agents (converged q-learners) to risk incurring high penalties. This is unlike the environment considered in this homework, where the optimal path is in a sense "risk-free", allowing Q-learning to outperform SARSA.

Homework: Maximization Bias

1. For the sake of clarity, we label the four outgoing actions from B as a_1, a_2, a_3 and a_4, from left to right, and say they belong to the action set A. For expected SARSA, we use the expected SARSA update rule to determine the state-action values:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \mathbb{E}_\pi \left[Q(S_{t+1}, A_{t+1}) | S_{t+1} \right] - Q(S_t, A_t) \right]$$

$$= Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_{a \in A} \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \quad (1)$$

Because all actions from B lead to a terminal state, we have that $Q(S_{t+1}, a) = 0$ for all $a \in A$ when $S_t = B$.

For a_1, on the first relevant sampled episode we have $R_{t+1} = 0$ giving:

$$Q(B, a_1) \leftarrow 0.7 + 0.2 \left[0 + 1 \times 4(\overset{0}{\cancel{0.25 \times 0}}) - 0.7 \right]$$

$$= 0.7 + 0.2 \left[-0.7 \right]$$

$$= 0.56. \quad (2)$$

Figure 1: SARSA vs Q-learning average return over number of episodes for the Windy Gridworld environment.

And on the next relevant sampled episode we get the same reward, giving:

$$Q(B, a_1) \leftarrow 0.56 + 0.2 \left[0 + 1 \times 4(0.25 \times 0)^{\,0} - 0.56 \right]$$
$$= 0.56 + 0.2 \left[-0.56 \right]$$
$$= 0.448. \tag{3}$$

For a_2, on the first relevant sampled episode we have $R_{t+1} = 1$, giving:

$$Q(B, a_2) \leftarrow 0.7 + 0.2 \left[1 + 1 \times 4(0.25 \times 0)^{\,0} - 0.7 \right]$$
$$= 0.7 + 0.2 \left[0.3 \right]$$
$$= 0.76. \tag{4}$$

And on the next relevant sampled episode, we get the same reward, giving:

$$Q(B, a_2) \leftarrow 0.76 + 0.2 \left[1 + 1 \times 4(0.25 \times 0)^{\,0} - 0.76 \right]$$
$$= 0.76 + 0.2 \left[0.24 \right]$$
$$= 0.808. \tag{5}$$

For a_3, on the first relevant sampled episode we have $R_{t+1} = 1$, which we know from the first update to a_2 gives us

$$Q(B, a_3) \leftarrow 0.76. \tag{6}$$

On the next relevant sampled episode, we have $R_{t+1} = 0$, giving:

$$\begin{aligned}
Q(B, a_2) &\leftarrow 0.76 + 0.2 \left[0 + 1 \times 4(\cancel{0.25 \times 0})^{0} - 0.76 \right] \\
&= 0.76 + 0.2 \left[-0.76 \right] \\
&= 0.608.
\end{aligned} \tag{7}$$

For a_4, on the first relevant sampled episode we have $R_{t+1} = 0$, which we know from the first update to a_1 gives us

$$Q(B, a_4) \leftarrow 0.56. \tag{8}$$

On the next relevant sampled episode, we have $R_{t+1} = 1$, giving:

$$\begin{aligned}
Q(B, a_1) &\leftarrow 0.56 + 0.2 \left[1 + 1 \times 4(\cancel{0.25 \times 0})^{0} - 0.56 \right] \\
&= 0.56 + 0.2 \left[0.44 \right] \\
&= 0.648.
\end{aligned} \tag{9}$$

For Q-learning, we use the Q-learning update rule to determine the state-action values:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \tag{10}$$

Note once again that since when $S_t = B$, S_{t+1} is always a terminal state, then like before $Q(S_{t+1}, a) = 0$ for all $a \in A$. Therefore, in this case, equation (10) reduces like equation (1) to

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} - Q(S_t, A_t) \right]. \tag{11}$$

Therefore, all the state-action values in state B are the same in Q-learning as for expected SARSA. For a clearer summary, refer to Table 1.

Table 1: Expected SARSA and Q-learning state-action pair values for the four available actions at state B after sampling two episodes per action.

	$Q(B, a_1)$	$Q(B, a_2)$	$Q(B, a_3)$	$Q(B, a_4)$
expected SARSA	0.448	0.808	0.608	0.648
Q-learning	0.448	0.808	0.608	0.648

2. To determine what the new $Q(A, L)$ value is when executing L in A after the 10 episodes, assuming that $Q(A, L)$ is still at 0.7, we use the same update rules as stated before, i.e. equation (1) for expected SARSA and equation (10) for Q-learning. Since

taking L at A leads to a terminal state, equations (1) and (10) once again reduce to equation (11). For both expected SARSA and Q-learning we therefore have:

$$\begin{aligned} Q(A, L) &\leftarrow 0.7 + 0.2 \left[0.7 - 0.7 \right] \\ &= 0.7 + 0.2 \left[0 \right] \\ &= 0.7. \end{aligned} \tag{12}$$

We apply the same process to determine what the new $Q(A, R)$ value is when executing R in A after the 10 episodes, assuming that $Q(A, R)$ is still at 0.7. However, the reduction to equation (11) is not possible in this case, since R from A does not transition to a terminal state. With expected SARSA we have

$$\begin{aligned} Q(A, R) &\leftarrow 0.7 + 0.2 \left[0 + 0.25 \left(0.448 + 0.808 + 0.608 + 0.648 \right) - 0.7 \right] \\ &= 0.6856. \end{aligned} \tag{13}$$

With Q-learning, we have

$$\begin{aligned} Q(A, R) &\leftarrow 0.7 + 0.2 \left[0 + 0.808 - 0.7 \right] \\ &= 0.7216. \end{aligned} \tag{14}$$

For a clearer summary, please refer to Table 2.

Table 2: Expected SARSA and Q-learning state-action pair values at A when executing R and L from A after the 10 sampled episodes.

	Expected SARSA	Q-learning
$Q(A, L)$	0.7	0.7
$Q(A, R)$	0.6856	0.7216

3. Assuming convergence to optimality for both Q-learning and Expected SARSA, we can obtain the true state-action values by utilising the Bellman optimality equation:

$$q_*(s, a) = \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \max_{a'} q_*(s', a') \right]. \tag{15}$$

The results of applying this equation to our MDP are summarised in Table 3.

Table 3: True state-action values after Expected SARSA and Q-learning convergence.

$Q_*(A, L)$	$Q_*(A, R)$	$Q_*(B, a_1)$	$Q_*(B, a_2)$	$Q_*(B, a_3)$	$Q_*(B, a_4)$
0.7	0.5	0.5	0.5	0.5	0.5

4. Maximization bias can be observed in the value of $Q(A, R)$ under Q-learning. We observe maximization bias here as the estimated value (reported in Table 2) are higher than the true values (reported in Table 3), i.e. they are positively biased. In general, both Q-learning and expected SARSA are affected by this bias, as both algorithms rely on a greedy (target) policy which requires the use of a maximum operator. When coupled with stochastic transitions and rewards, this generally leads to a positive bias in the estimated state-action values. In our particular example, we only observe maximization bias under Q-learning since here because we see an (erroneous) positive increase in value from the initial value, from 0.7 to 0.7216, while with Expected SARSA the value is (correctly) decreased from 0.7 to 0.6856.

Table 4: State-action values at state B due to Double Q-learning after the 10 observed episodes, alternating updates to Q_1 and Q_2

	B, a_1	B, a_2	B, a_3	B, a_4
Q_1	0.56	0.76	0.76	0.56
Q_2	0.56	0.76	0.56	0.76

5. Double Q-learning circumvents the issue of maximization bias in Q-learning by using two independent estimates, Q_1 and Q_2, of the true value function q. This is unlike vanilla Q-learning where we use a single estimate Q. The two estimates afford us the possibility of using one estimate for determining the greedy action $A^* = \arg\max_a Q_1(a)$ and the other for estimating its value $Q_2(A^*) = Q_2(\arg\max_a Q_1(a))$. The latter estimate is then unbiased: $\mathbb{E}[Q_2(A^*)] = q(A^*)$. We can then repeat the process with Q_1 and Q_2 swapped to obtain another unbiased estimate. More specifically, in double Q-learning for any given timestep t we would, with probability 0.5, use the following update rule:

$$Q_1(S_t, A_t) \leftarrow Q_1(A_t, A_t) + \alpha \left[R_t + \gamma Q_2(S_{t+1}, \arg\max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right],$$
(16)

and flip the roles of Q_1 and Q_2 otherwise. For a more concrete example, consider the estimate of $q(A, R)$, whose true value is 0.5 but we have estimated to be 0.7216 in vanilla Q-learning as shown in Table 2. To estimate this under Double Q-learning, we repeat the work of question 1., using Double Q-learning rather than vanilla Q-learning, to obtain $Q_1(B, a), Q_2(B, a) \; \forall a \in A$. For simplicity, we deterministically alternate updates to Q_1 and Q_2 between the observed rewards, still assuming an initialization of 0.7. The results are summarised in Table 4. Using the values in Table 4, we can then continue similarly to question 2., assuming that $Q_1(A, R)$ and $Q_2(A, R)$ stayed at 0.7 after the first 10 iterations. We apply Double Q-learning breaking ties arbitrarily and obtain:

$$Q_1(A, R) \leftarrow 0.7 + 0.2 \left[0 + Q_2(B, \arg\max_a Q_1(B, a)) - 0.7 \right]$$
$$= 0.7 + 0.2 \left[Q_2(B, a_3) - 0.7 \right]$$
$$= 0.7 + 0.2 \left[0.56 - 0.7 \right]$$
$$= 0.672.$$
(17)

We can repeat this process with Q_1 and Q_2 swapped to obtain another unbiased estimate:

$$Q_2(A, R) \leftarrow 0.7 + 0.2 \left[0 + Q_1(B, \arg\max_a Q_2(B, a)) - 0.7 \right]$$
$$= 0.7 + 0.2 \left[Q_1(B, a_4) - 0.7 \right]$$
$$= 0.7 + 0.2 \left[0.56 - 0.7 \right]$$
$$= 0.672.$$
(18)

We see that our estimate has now been reduced from 0.7216 to 0.672, correctly updating with a decrease from 0.7 towards the true value of 0.5, hence circumventing maximization bias.

Homework: Gradient Descent Methods

1. The true value of a state $v_\pi(S_t)$ is defined to be the expected return at that state, i.e.

$$v_\pi(S_t) = \mathbb{E}\left[G_t \mid S_t = s\right]. \tag{19}$$

Since the Monte Carlo target *is* the return G_t, estimated using direct samples of the true value, it has the same expected value as $v_\pi(S_t)$ and is therefore an unbiased estimate of the true value of a state by definition.

2. The temporal difference (TD) error, δ_t is given by

$$\delta_t \doteq R_{t+1} + \gamma V(S_{t+1}) - V(S_t). \tag{20}$$

A weight update that minimizes the mean squared temporal difference error $\overline{\delta_t}$ can be derived by considering the definition of mean square *value* error $\overline{\text{VE}}$:

$$\overline{\text{VE}} \doteq \sum_{s \in \mathcal{S}} \mu(s) \left[v_\pi(s) - \hat{v}(s, \mathbf{w})\right]^2, \tag{21}$$

where $\hat{v}$ is our approximation and v_π is the true value. Here we have the squared error term in the brackets, weighted by "how much we care about it", $\mu(s)$ in a sum. We can replace the error term in the brackets with the TD error given in equation (20), obtaining

$$\overline{\delta_t} = \sum_{s_t, s_{t+1} \in \mathcal{S}} \mu(s) \left[R_{t+1} + \gamma \hat{v}(s_{t+1}, \mathbf{w}) - \hat{v}(s_t, \mathbf{w}))\right]^2, \tag{22}$$

where we are now summing over all possible state transitions (s_t, s_{t+1}). We can then find the weight update rule that minimizes the mean squared TD error by taking the gradient with respect to $\mathbf{w}$:

$$
\begin{aligned}
\mathbf{w}_{t+1} &\leftarrow \mathbf{w}_t - \frac{1}{2}\alpha \nabla \left[R_{t+1} + \gamma \hat{v}(s_{t+1}, \mathbf{w}) - \hat{v}(s_t, \mathbf{w}))\right]^2 \\
&= \mathbf{w}_t + \alpha \left[R_{t+1} + \gamma \hat{v}(s_{t+1}, \mathbf{w}) - \hat{v}(s_t, \mathbf{w}))\right] \nabla \left[\gamma \hat{v}(s_{t+1}, \mathbf{w}) - \hat{v}(s_t, \mathbf{w})\right] \\
&= \mathbf{w}_t + \alpha \delta_t \nabla \left[\underline{\gamma \hat{v}(s_{t+1}\mathbf{w}) - \hat{v}(s_t, \mathbf{w})}\right]
\end{aligned} \tag{23}
$$

We note that we no longer make use of information from the true target value, v_π, instead only accounting for changes on the estimates. We further note that unlike Semi-Gradient TD, we have an extra term (underlined), necessary to compute the gradient correctly. In Semi-Gradient TD, this term is missing and in this sense, only a portion of the true gradient is considered, hence the name "semi-gradient".

3. Semi-gradient methods indeed have some advantages over Monte Carlo approaches. A bootstrapping method would be preferred over a Monte Carlo approach in the case of the Mountain Car problem because it would enable continuous and online learning. In simpler words: non-stop learning that can occur interleaved with the agent actions. This is desired in this case because a given episode terminates only when the car reaches the goal, making learning much slower if relying on Monte Carlo methods, which require the use of the return.