

Exercise Set 3 - Reinforcement Learning

Advanced TD methods and approximation

Giulio Starace - 13010840

September 24, 2022

Homework: Coding Assignment - Temporal Difference Learning

1. Coding answers have been submitted on codegra under the group “stalwart cocky sawly”.
2. Hello World

Homework: Maximization Bias

1. For the sake of clarity, we label the four outgoing actions from B as a_1, a_2, a_3 and a_4 , from left to right, and say they belong to the action set A . For expected SARSA, we use the expected SARSA update rule to determine the state-action values:

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \mathbb{E}_\pi [Q(S_{t+1}, A_{t+1}) | S_{t+1}] - Q(S_t, A_t)] \\ &= Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_{a \in A} \pi(a | S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \end{aligned} \quad (1)$$

Because all actions from B lead to a terminal state, we have that $Q(S_{t+1}, a) = 0$ for all $a \in A$ when $S_t = B$.

For a_1 , on the first relevant sampled episode we have $R_{t+1} = 0$ giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.7 + 0.2 \left[0 + 1 \times 4(0.25 \times 0) - 0.7 \right] \\ &= 0.7 + 0.2 [-0.7] \\ &= 0.56. \end{aligned} \quad (2)$$

And on the next relevant sampled episode we get the same reward, giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.56 + 0.2 \left[0 + 1 \times 4(0.25 \times 0) - 0.56 \right] \\ &= 0.56 + 0.2 [-0.56] \\ &= 0.448. \end{aligned} \quad (3)$$

For a_2 , on the first relevant sampled episode we have $R_{t+1} = 1$, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.7 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.7 \right] \\ &= 0.7 + 0.2 [0.3] \\ &= 0.76. \end{aligned} \tag{4}$$

And on the next relevant sampled episode, we get the same reward, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.76 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.76 \right] \\ &= 0.76 + 0.2 [0.24] \\ &= 0.808. \end{aligned} \tag{5}$$

For a_3 , on the first relevant sampled episode we have $R_{t+1} = 1$, which we know from the first update to a_2 gives us

$$Q(B, a_3) \leftarrow 0.76. \tag{6}$$

On the next relevant sampled episode, we have $R_{t+1} = 0$, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.76 + 0.2 \left[0 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.76 \right] \\ &= 0.76 + 0.2 [-0.76] \\ &= 0.608. \end{aligned} \tag{7}$$

For a_4 , on the first relevant sampled episode we have $R_{t+1} = 0$, which we know from the first update to a_1 gives us

$$Q(B, a_4) \leftarrow 0.56. \tag{8}$$

On the next relevant sampled episode, we have $R_{t+1} = 1$, giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.56 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.56 \right] \\ &= 0.56 + 0.2 [0.44] \\ &= 0.648. \end{aligned} \tag{9}$$

For Q-learning, we use the Q-learning update rule to determine the state-action values:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \tag{10}$$

Note once again that since when $S_t = B$, S_{t+1} is always a terminal state, then like before $Q(S_{t+1}, a) = 0$ for all $a \in A$. Therefore, in this case, equation (10) reduces like equation (1) to

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} - Q(S_t, A_t)]. \tag{11}$$

Therefore, all the state-action values in state B are the same in Q-learning as for expected SARSA. For a clearer summary, refer to Table 1.

Table 1: Expected SARSA and Q-learning state-action pair values for the four available actions at state B after sampling two episodes per action.

	$Q(B, a_1)$	$Q(B, a_2)$	$Q(B, a_3)$	$Q(B, a_4)$
expected SARSA	0.448	0.808	0.608	0.648
Q-learning	0.448	0.808	0.608	0.648

2. To determine what the new $Q(A, L)$ value is when executing L in A after the 10 episodes, assuming that $Q(A, L)$ is still at 0.7, we use the same update rules as stated before, i.e. equation (1) for expected SARSA and equation (10) for Q-learning. Since taking L at A leads to a terminal state, equations (1) and (10) once again reduce to equation (11). For both expected SARSA and Q-learning we therefore have:

$$\begin{aligned}
 Q(A, L) &\leftarrow 0.7 + 0.2 [0.7 - 0.7] \\
 &= 0.7 + 0.2 [0] \\
 &= 0.7.
 \end{aligned} \tag{12}$$

We apply the same process to determine what the new $Q(A, R)$ value is when executing R in A after the 10 episodes, assuming that $Q(A, R)$ is still at 0.7. However, the reduction to equation (11) is not possible in this case, since R from A does not transition to a terminal state. With expected SARSA we have

$$\begin{aligned}
 Q(A, R) &\leftarrow 0.7 + 0.2 [0 + 0.25 (0.448 + 0.808 + 0.608 + 0.648) - 0.7] \\
 &= 0.6856.
 \end{aligned} \tag{13}$$

With Q-learning, we have

$$\begin{aligned}
 Q(A, R) &\leftarrow 0.7 + 0.2 [0 + 0.808 - 0.7] \\
 &= 0.7216.
 \end{aligned} \tag{14}$$

For a clearer summary, please refer to Table 2.

Table 2: Expected SARSA and Q-learning state-action pair values at A when executing R and L from A after the 10 sampled episodes.

	Expected SARSA	Q-learning
$Q(A, L)$	0.7	0.7
$Q(A, R)$	0.6856	0.7216

3. Assuming convergence to optimality for both Q-learning and Expected SARSA, we can obtain the true state-action values by utilising the Bellman optimality equation:

$$q_*(s, a) = \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \max_{a'} q_*(s', a') \right]. \tag{15}$$

The results of applying this equation to our MDP are summarised in Table 3.

4. Maximization bias can be observed in all estimated state-action values reported in Tables 1 and 2, except for $Q(A, L)$ and $Q(B, a_1)$. We observe maximization bias

Table 3: True state-action values after Expected SARSA and Q-learning convergence.

$Q_*(A, L)$	$Q_*(A, R)$	$Q_*(B, a_1)$	$Q_*(B, a_2)$	$Q_*(B, a_3)$	$Q_*(B, a_4)$
0.7	0.5	0.5	0.5	0.5	0.5

here as the estimated values are higher than the true values reported in Table 3, i.e. they are positively biased. Both Q-learning and expected SARSA are affected by this bias, as both algorithms rely on a greedy (target) policy which requires the use of a maximum operator. When coupled with stochastic transitions and rewards, such as in our MDP from state B , this generally leads to a positive bias in the estimated state-action values.

5. Double Q-learning circumvents the issue of maximization bias in Q-learning by using two independent estimates, Q_1 and Q_2 , of the true value function q . This is unlike vanilla Q-learning where we use a single estimate Q . The two estimates afford us the possibility of using one estimate for determining the greedy action $A^* = \arg \max_a Q_1(a)$ and the other for estimating its value $Q_2(A^*) = Q_2(\arg \max_a Q_1(a))$. The latter estimate is then unbiased: $\mathbb{E}[Q_2(A^*)] = q(A^*)$. We can then repeat the process with Q_1 and Q_2 swapped to obtain another unbiased estimate. More specifically, in double Q-learning for any given timestep t we would, with probability 0.5, use the following update rule:

$$Q_1(S_t, A_t) \leftarrow Q_1(S_t, A_t) + \alpha \left[R_t + \gamma Q_2(S_{t+1}, \arg \max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right], \quad (16)$$

and flip the roles of Q_1 and Q_2 otherwise. For a more concrete example, consider the estimate of $q(A, R)$, whose true value is 0.7 but we have estimated to be 0.7216 in vanilla Q-learning as shown in Table 2. Under Double Q-learning, still assuming an initialization of 0.7, our estimate would now look like:

$$\begin{aligned} Q_1(a, r) &\leftarrow 0.7 + 0.2 [0 + Q_2(b, a_2) - 0.7] \\ &= 0.7 + 0.2 [0.7 - 0.7] \\ &= 0.7. \end{aligned} \quad (17)$$

We can repeat this process with Q_1 and Q_2 swapped to obtain another unbiased estimate:

$$\begin{aligned} Q_2(A, R) &\leftarrow 0.7 + 0.2 [0 + Q_1(B, a_2) - 0.7] \\ &= 0.7 + 0.2 [0.7 - 0.7] \\ &= 0.7. \end{aligned} \quad (18)$$

We see that our estimate has now been reduced from 0.7216 to 0.7, the true value, thus circumventing the maximization bias issue.

Homework: Gradient Descent Methods

1. The true value of a state $v_\pi(S_t)$ is defined to be the expected return at that state, i.e.

$$v_\pi(S_t) = \mathbb{E}[G_t \mid S_t = s]. \quad (19)$$

Since the Monte Carlo target *is* the return G_t , estimated using direct samples of the true value, it has the same expected value as $v_\pi(S_t)$ and is therefore an unbiased estimate of the true value of a state by definition.

2. hello world

3. hello world