

Exercise Set 3 - Reinforcement Learning

Advanced TD methods and approximation

Giulio Starace - 13010840

September 24, 2022

Homework: Coding Assignment - Temporal Difference Learning

1. Coding answers have been submitted on codegra under the group “stalwart cocky sawly”.
2. Hello World

Homework: Maximization Bias

1. For the sake of clarity, we label the four outgoing actions from B as a_1, a_2, a_3 and a_4 , from left to right, and say they belong to the action set A . For expected SARSA, we use the expected SARSA update rule to determine the state-action values:

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \mathbb{E}_\pi [Q(S_{t+1}, A_{t+1}) | S_{t+1}] - Q(S_t, A_t)] \\ &= Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_{a \in A} \pi(a | S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \end{aligned} \quad (1)$$

Because all actions from B lead to a terminal state, we have that $Q(S_{t+1}, a) = 0$ for all $a \in A$ when $S_t = B$.

For a_1 , on the first relevant sampled episode we have $R_{t+1} = 0$ giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.7 + 0.2 \left[0 + 1 \times 4(0.25 \times 0) - 0.7 \right] \\ &= 0.7 + 0.2 [-0.7] \\ &= 0.56. \end{aligned} \quad (2)$$

And on the next relevant sampled episode we get the same reward, giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.56 + 0.2 \left[0 + 1 \times 4(0.25 \times 0) - 0.56 \right] \\ &= 0.56 + 0.2 [-0.56] \\ &= 0.448. \end{aligned} \quad (3)$$

For a_2 , on the first relevant sampled episode we have $R_{t+1} = 1$, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.7 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.7 \right] \\ &= 0.7 + 0.2 [0.3] \\ &= 0.76. \end{aligned} \tag{4}$$

And on the next relevant sampled episode, we get the same reward, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.76 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.76 \right] \\ &= 0.76 + 0.2 [0.24] \\ &= 0.808. \end{aligned} \tag{5}$$

For a_3 , on the first relevant sampled episode we have $R_{t+1} = 1$, which we know from the first update to a_2 gives us

$$Q(B, a_3) \leftarrow 0.76. \tag{6}$$

On the next relevant sampled episode, we have $R_{t+1} = 0$, giving:

$$\begin{aligned} Q(B, a_2) &\leftarrow 0.76 + 0.2 \left[0 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.76 \right] \\ &= 0.76 + 0.2 [-0.76] \\ &= 0.608. \end{aligned} \tag{7}$$

For a_4 , on the first relevant sampled episode we have $R_{t+1} = 0$, which we know from the first update to a_1 gives us

$$Q(B, a_4) \leftarrow 0.56. \tag{8}$$

On the next relevant sampled episode, we have $R_{t+1} = 1$, giving:

$$\begin{aligned} Q(B, a_1) &\leftarrow 0.56 + 0.2 \left[1 + 1 \times 4 \overset{0}{\cancel{(0.25 \times 0)}} - 0.56 \right] \\ &= 0.56 + 0.2 [0.44] \\ &= 0.648. \end{aligned} \tag{9}$$

For Q-learning, we use the Q-learning update rule to determine the state-action values:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_{a \in A} Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \tag{10}$$

Note once again that since when $S_t = B$, S_{t+1} is always a terminal state, then like before $Q(S_{t+1}, a) = 0$ for all $a \in A$. Therefore, in this case, equation (10) reduces like equation (1) to

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} - Q(S_t, A_t)]. \tag{11}$$

Therefore, all the state-action values in state B are the same in Q-learning as for expected SARSA. For a clearer summary, refer to Table 1.

Table 1: Expected SARSA and Q-learning state-action pair values for the four available actions at state B after sampling two episodes per action.

	$Q(B, a_1)$	$Q(B, a_2)$	$Q(B, a_3)$	$Q(B, a_4)$
expected SARSA	0.448	0.808	0.608	0.648
Q-learning	0.448	0.808	0.608	0.648

2. To determine what the new $Q(A, L)$ value is when executing L in A after the 10 episodes, assuming that $Q(A, L)$ is still at 0.7, we use the same update rules as stated before, i.e. equation (1) for expected SARSA and equation (10) for Q-learning. Since taking L at A leads to a terminal state, equations (1) and (10) once again reduce to equation (11). For both expected SARSA and Q-learning we therefore have:

$$\begin{aligned}
 Q(A, L) &\leftarrow 0.7 + 0.2 [0.7 - 0.7] \\
 &= 0.7 + 0.2 [0] \\
 &= 0.7.
 \end{aligned} \tag{12}$$

We apply the same process to determine what the new $Q(A, R)$ value is when executing R in A after the 10 episodes, assuming that $Q(A, R)$ is still at 0.7. However, the reduction to equation (11) is not possible in this case, since R from A does not transition to a terminal state. With expected SARSA we have

$$\begin{aligned}
 Q(A, R) &\leftarrow 0.7 + 0.2 [0 + 0.25 (0.448 + 0.808 + 0.608 + 0.648) - 0.7] \\
 &= 0.6856.
 \end{aligned} \tag{13}$$

With Q-learning, we have

$$\begin{aligned}
 Q(A, R) &\leftarrow 0.7 + 0.2 [0 + 0.808 - 0.7] \\
 &= 0.7216.
 \end{aligned} \tag{14}$$

For a clearer summary, please refer to Table 2.

Table 2: Expected SARSA and Q-learning state-action pair values at A when executing R and L from A after the 10 sampled episodes.

	Expected SARSA	Q-learning
$Q(A, L)$	0.7	0.7
$Q(A, R)$	0.6856	0.7216

3. hello world
4. hello world
5. hello world

Homework: Gradient Descent Methods

1. hello world
2. hello world
3. hello world